

Hee Won Lee

36 Liberty Ridge Rd, Basking Ridge, NJ 07920, USA (*Green Card Holder*)
knowpd@gmail.com • +1 (919) 800-8993 • <http://knowpd.github.com>

EDUCATION

North Carolina State University, Raleigh, North Carolina, USA

- Ph.D. in Computer Science Aug 2009 – May 2015
 - Thesis: Network Emulation with Adaptive Time Dilation
 - Adviser: Prof. Mihail L. Sichitiu, Co-adviser: Prof. David Thunte
 - Focus: Networked Systems, Distributed Computing, Discrete Simulation.
 - Cumulative GPA: 4.0 / 4.0

Carnegie Mellon University, Pittsburgh, Pennsylvania, USA

- M.E. in Carnegie Mellon University Aug 2004 – Aug 2005
 - Cumulative GPA: 3.4 / 4.0

Korea University, Seoul, South Korea

- B.E. in Electrical Engineering Mar 1995 – Feb 2002
 - Cumulative GPA: 3.3 / 4.0

PROFESSIONAL EXPERIENCES

AT&T Labs Research, Bedminster, New Jersey, USA

- Principal Inventive Scientist Mar 2015 – Jul 2020
 - 2019–2020
 - Redis with Kubernetes
 - Evaluated the performance of in-memory databases (Redis vs. Aerospike)
 - Designed and implemented Redis High Availability with Kubernetes in CI/CD environment
 - Integrated Redis HA helm charts into Radio Access Network Intelligent Controller release 4 (field trial version)
 - Persistent Memory for In-memory Database & Content Delivery Network (CDN)
 - Evaluated the performance of Redis with Intel Optane Persistent Memory
 - Designed a new key-value store architecture for Persistent Memory
 - Collaborated with Intel team to analyze the performance characteristics
 - Analyzed cost benefits derived from applying Intel Optane Persistent Memory to AT&T's vCDN
 - 2017–2018
 - Software-Defined Storage
 - Designed and implemented Ceph helm charts for Kubernetes with an one-osd-per-pod approach
 - Performed extensive resiliency tests for containerized Ceph storage with Kubernetes
 - Developed best recovery practices from various failure scenarios for containerized Ceph storage with Kubernetes
 - Software-Defined Caching
 - Evaluated the performance of page cache & dm-cache in virtualized environments
 - Characterized the performance of multi-tier caching in Ceph storage (host/vm page cache, rbd cache, ceph cache-tier)
 - Advanced storage networking for ultra low latency with RDMA
 - Set up a testbed for NVMe over Fabrics with CPU offloading on RoCE-based lossless networks
 - Evaluated the performance and verified zero CPU usage on the target host machine
 - 2015–2016
 - Multi-site Network Traffic Shaper for Quality of Service
 - Designed and implemented a system that controls network bandwidth/delay between multi-site OpenStack clusters
 - Open sourced the code at <https://github.com/att/netarbitrator/tree/master/multisite-netemu>

AT&T Labs Research, Bedminster, New Jersey, USA

- Summer Intern May 2014 – Aug 2014
 - Project: I designed and implemented I/O traffic protection mechanisms for cloud storage
 - Mentor: Moo-Ryong Ra, PhD

North Carolina State University, Raleigh, North Carolina, USA

- Research Assistant Jan 2010 – Dec 2014
 - Projects:
 - I designed and implemented a network emulation framework for evaluating distributed applications running on diverse OSs (Linux, FreeBSD, Windows, JunOS)
 - I modified the QEMU-KVM hypervisor to synchronize distributed VMs with virtual time
 - Advisors: Prof. Mihail L. Sichitiu and Prof. David Thunte

KT (Korea Telecom), Daejeon, South Korea

- Member of Technical Staff Jan 2002 – Aug 2013

- Projects:
 - I designed and implemented the systems that monitor and control Internet Backbone and VoIP networks.
 - I performed network capacity planning for Internet Backbone and IPTV Services with OPNET Network simulator.

PUBLICATIONS

CONFERENCES

- [6] **EF-Dedup: Enabling Collaborative Data Deduplication at the Network Edge.**
Shijing Li, Tian Lan, Bharath Balasubramanian, Moo-ryong Ra, Hee Won Lee, and Panta Rajesh.
In *IEEE 39th International Conference on Distributed Computing Systems (ICDCS 2019)*.
Dallas, TX, USA, Jul 2019.
- [5] **Fighting with Unknowns: Estimating the Performance of Scalable Distributed Storage Systems with Minimal Measurement Data.**
Moo-ryong Ra and Hee Won Lee.
In *IEEE 35th Symposium on Mass Storage Systems and Technologies (MSST 2019)*.
Santa Clara, CA, USA, May 2019.
- [4] **MIST: Mitigating Host-Side Interference for Storage Traffic in Virtualized Data Centers.**
Hee Won Lee and Moo-ryong Ra
In *IEEE 9th International Conference on Cloud Computing (CLOUD 2016)*.
San Francisco, CA, USA, Jun 2016.
- [3] **Integrated Simulation and Emulation using Adaptive Time Dilation.**
Hee Won Lee, David Thuente, and Mihail L. Sichitiu.
In *ACM 2nd SIGSIM Conference on Principles of Advanced Discrete Simulation (PADS 2014)*.
Santa Clara, CA, USA, May 2014.
- [2] **A Resource Management System for Next Generation Services.**
Seung-Hee Han, Bom-Su Kim, Chan-Kyou Hwang, Hee Won Lee, Byung-deok Chung.
In *IEEE International Conference on Advanced Technologies for Communications 2019*.
Hai Phong, Vietnam, Oct 2009.
- [1] **Reusability Enhancement by Using Flexible Topology Architecture for Network Management System.**
Hee Won Lee, Chan Kyou Hwang, Jae-Hyoung Yoo, Ho-Jin Choi, Sungwon Kang, and Dan H. Lee.
In *IEEE/ACIS 7th International Conference on Computer and Information Science (ICIS 2008)*.
Portland, Oregon, USA, May 2008.

JOURNALS

- [2] **Network Link Emulation With Adaptive Time Dilation.**
Hee Won Lee, Mihail L. Sichitiu, and David Thuente.
Journal of Parallel and Distributed Computing, vol. 104, pp. 88–98.
Jun 2017.
- [1] **High-performance Emulation of Heterogeneous Systems using Adaptive Time Dilation.**
Hee Won Lee, Mihail L. Sichitiu, and David Thuente.
International Journal of High Performance Computing Applications, vol. 29, issue 2.
May 2015.

WORKSHOPS, POSTERS AND TECHNICAL REPORTS

- [3] **Accelerating Applications in the Fast-Moving Devices with Proactive Provisioning (Poster).**
HyunJong Lee, Hee Won Lee, Moo-Ryong Ra, Yu Xiang, and Jason Flinn.
In *Proceedings of the 17th Annual International Conference on Mobile Systems, Applications, and Services (MobiSys 2019)*.
Seoul, South Korea, Jun 2019.
- [2] **IOArbiter: Dynamic Provisioning of Backend Block Storage in the Cloud.**
Moo-Ryong Ra and Hee Won Lee.
arXiv preprint arXiv:1904.09984.
Apr 2019.

- [1] **End-User IPTV Traffic Measurement of Residential Broadband Access Networks.**
 Young J. Won, Mi-Jung Choi, Byung-Chul Park, James W. Hong, Hee Won Lee, Chan Kyu Hwang,
 and Jae-Hyoung Yoo.
 In *IEEE Network Operations and Management Symposium (NOMS) Workshops*.
 Apr 2008.

PATENTS

ISSUED

- [1] **Dynamic Provisioning of Storage in the Cloud.**
 Moo-ryong Ra and Hee Won Lee.
 US 10,530,703 B2.
 Date of Patent: Jan 7, 2020.

SOFTWARE SYSTEM EXPERIENCES

Memory & Storage

- Redis In-memory Database 2019 – 2020
 - Designed a framework that converts volatile indexes into their persistent counterparts
 - Designed and implemented Redis high availability with Kubernetes helm charts
- Ceph Containerization 2017
 - Designed and implemented containerized Ceph with Kubernetes helm charts
- Software Defined Caching 2017
 - Evaluated the performance of page cache and dm-cache in a virtual environment
 - Evaluated the performance of Ceph with dm-cache
 - Designed a system architecture for providing coordinated caching layers in a multi-tenancy environment
- All Flash Array Storage System 2016 – 2017
 - Evaluated and compared the performance of SATA SSD and NVMe SSD
 - Set up storage networking using RDMA over Converged Ethernet (RoCE)
 - Designed a system architecture for making a commodity SSD array a usable iSCSI block storage backend for an OpenStack cloud
- Backend QoS for Cloud Block Storage 2015
 - Designed and implemented the backend of OpenStack Cinder for dynamic configuration and QoS support
 - Containerized OpenStack's cinder-volume
- I/O Traffic Protection for Cloud Storage [Python] 2014
 - Designed and implemented a storage I/O bandwidth reservation algorithm that guarantees the minimum I/O bandwidth of a virtual host which is remotely connected by a storage area network (SAN) to a shared storage volume through iSCSI
 - Designed a NUMA-aware CPU pinning algorithm that protects I/O traffic from CPU interference
 - Evaluated the system using diverse types of I/O pattern with FIO and realistic workloads with Filebench
- Multi-Threaded Text File Indexing [C++] 2014
 - Created a thread for searching text file paths and multiple worker threads for counting words using Boost Filesystem and Thread Libraries
 - Implemented multiple producer, multiple consumer thread-safe queue and map
- Memory Control [Python] 2014
 - Designed and implemented a Cgroup-based memory ceiling algorithm to protect I/O traffic bandwidth in overprovisioned cloud environments
- Demand Paging [C] 2013
 - Implemented a demand paging system that allows for more address space than physically available one (XINU)

Networking

- Multi-site Network Emulation [Ansible] 2016
 - Designed and implemented a multi-site network emulation system that allows for stitching multiple OpenStack clusters and emulating the network latency/ bandwidth
 - Open sourced: <https://github.com/att/netarbiter>
- Network Link Emulation [C++, Python] 2014
 - Interconnected distributed VMs through virtual links using Tc/NetEm with adaptive time dilation
- Software-Defined Networking [Python] 2012
 - Used OpenFlow POX controller and Open vSwitch to migrate VirtualBox-based virtual networks
- Routing Algorithms [Java] 2012
 - Implemented modified Dijkstra's algorithm with negative weights
 - Implemented vertex/edge disjoint paths and elementary circuits search algorithms
 - Implemented CPLEX code with Integer Linear Programming for Routing and Wavelength Assignment

- **Queueing Systems [Java]** 2010
 - Implemented discrete event simulation models for queueing systems (M/M/1, M/M/m, M/G/m)
- **Wireless Networking [C++]** 2010
 - Designed and implemented an OMNeT++ model for IVG (Inter-Vehicle Geocast) routing protocol in vehicle-to-vehicle ad-hoc network simulation
 - Connected KVM-based virtual nodes through a wireless simulation model of OMNeT++
- **Network Management Systems [C]** 2002 – 2008
 - Designed and developed the systems that monitor and control Internet Backbone and VoIP networks using Net-SNMP (Small Network Management Protocol) and Oracle database

Virtual Machine

- **Hybrid Simulation and Emulation System [C++, Python]** 2014
 - Built an integrated simulation (NS-3) and emulation (QEMU-KVM) framework in a distributed environment
- **High-Performance Emulation System [C++, Python]** 2010 – 2013
 - Emulated higher performance with a time dilation technique using unmodified OSs (Linux, FreeBSD, Windows, and Junos) and real application workloads (VLC media player)
 - Modified the QEMU-KVM hypervisor for virtual time Peer-to-Peer File Sharing System [Java] 2009
 - Mapped a file signature (created by encoding text document with m-bit string) onto a hash space
 - Used Apache Mina, Apache FtpServer, and JavaDB
- **Virtual Time Synchronization [C++]** 2014
 - Designed and implemented a spinlock on shared memory to synchronize distributed VMs (processes created by KVM hypervisor) and simulation nodes (processes created by NS-3) with virtual time
 - Implemented a virtual time synchronization daemon with TCP/UDP sockets using Boost Asio Library

CPU

- **CPU Share Controller [Python]** 2014
 - Designed and implemented a CPU control algorithm that enforces CPU ceiling and allocates a relative share of CPU time using Linux Control Groups (Cgroups)
- **Prioritized Preemptive Schedulers [C]** 2013
 - Implemented a prioritized preemptive scheduler using POSIX thread library
 - Designed and implemented a priority-based scheduling algorithm using XINU kernel (a small UNIX OS)

RESEARCH MENTORING

- **Simran Singh**, North Carolina State University
 - Project: UAV real-time video applications with cellular infrastructure support
 - 2019 fall internship at AT&T
- **Madhava Krishnan**, Virginia Tech
 - Project: A general purpose in-memory key-value store architecture for emerging persistent memory
 - 2019 summer internship at AT&T
- **HyunJong (Joseph) Lee**, University of Michigan–Ann Arbor
 - Project: Predictive migration of edge computations
 - 2018 summer internship at AT&T
- **Zhang Liu**, University of Colorado–Boulder
 - Project: Multi-tier cache orchestration with miss ratio curve
 - 2017 summer internship at AT&T

TEACHING EXPERIENCES

North Carolina State University, Raleigh, North Carolina, USA

- **Outstanding Teaching Assistant**
 - For excellent service as a teaching assistant during the 2010-11 calendar year, the Department of Computer Science recognizes Hee Won Lee as an Outstanding Teaching Assistant
 - Awarded the day of 29 April 2011
- **Teaching Assistant**
 - Operating System (2010, 2012, 2013)
 - Algorithm (2011)
 - Computer Networks (2011, 2012)
 - Internet Protocols (2014)
 - Software Engineering (2010)
- **Guest lecturer for Computer Networks (CSC401)**
 - Lectured on Switching and Bridging on February 26, 2013

REFERENCES

- **Dr. Yih-Farn (Robin) Chen**
(1987–2020 at AT&T)
Director Inventive Science
AT&T Labs Research
1 AT&T Way, Bedminster, NJ 07921, USA
rccym@gmail.com • +1 (973) 960-1594
- **Prof. Mihail L. Sichitiu**
Department of Electrical and Computer Engineering
North Carolina State University
890 Oval Drive, 3114 Engineering Building II, Raleigh, NC 27606, USA
mlsichit@ncsu.edu • +1 (919) 515-7348
- **Dr. Moo-Ryong Ra**
(Previously, Principal Inventive Scientist at AT&T Labs Research during 2013–2019)
Senior Software Engineer
Amazon
130 Lytton Ave, Palo Alto, CA 94301, USA
mooryor@amazon.com • +1 (240) 583-0968
- **Prof. Sangtae Ha**
Department of Computer Science
University of Colorado Boulder
ECCR 1B14, 1045 Regent Drive 430 UCB, Boulder CO, 80309, USA
sangtae.ha@colorado.edu • +1 (303) 492-7031

[CV compiled on 2020-07-24]